

🔍 Problem Analysis

Problem Summary:

The AzureOpenAiChatModel has a race condition in streaming responses where chat chunks can arrive out of order due to Flux::flatMap not providing ordering guarantees. This causes streamed content like 'data: <A>' followed by 'data: ' to potentially be received as ' => <A>'.

- Key Requirements:**
- Maintain strict ordering of streaming chat response chunks
 - Preserve existing Spring AI chat model API contract and behavior
 - Ensure compatibility with Azure OpenAI streaming protocol
 - Maintain thread-safety in reactive streaming operations
 - Support concurrent processing while preserving order
 - Avoid breaking changes to existing chat model implementations
 - Maintain performance characteristics of streaming responses
 - Handle tool execution flow without introducing new race conditions

- Design Decisions:**
- Replace Flux::flatMap with flatMapSequential to enforce ordering guarantees
 - Apply the ordering fix to both main chat response stream and tool execution flow
 - Keep the existing reactive pipeline structure intact to minimize risk
 - Maintain the same merge and processing logic for chat completions
 - Preserve all existing error handling and response mapping behaviors

- Solution Approach:**
- Primary approach: Replace flatMap with flatMapSequential to maintain emission order
 - Applied consistently across both chat response processing and tool execution flows
 - Minimal invasive change targeting only the specific ordering

🔄 Backport Assessment

✅ APPROVE

Classification: Bug Fix

Risk Level: Low - minimal code change addressing concurrency issue without affecting public interfaces

Reasoning:

This is a clear bug fix that resolves a reordering issue in Azure OpenAI chat streaming without changing any public APIs. The change is limited to internal reactive stream processing and maintains backward compatibility while fixing a concurrency-related bug.

Recommendations:

Safe for backporting. Consider testing the fix in 1.0.x environment to ensure the reactive streams behavior is consistent across Spring versions. No special migration considerations needed.